



# Nuvolaris Trainings

## Developing Kubernetes Operators in Python

Part 1: Kubernetes Operators

<https://www.nuvolaris.io>

# Agenda

- Kubernetes Introduction
- Development Environment
- Kubernetes Operators
- Creating an Operator
- Deploying an Operator
- Contributing to Nuvolaris

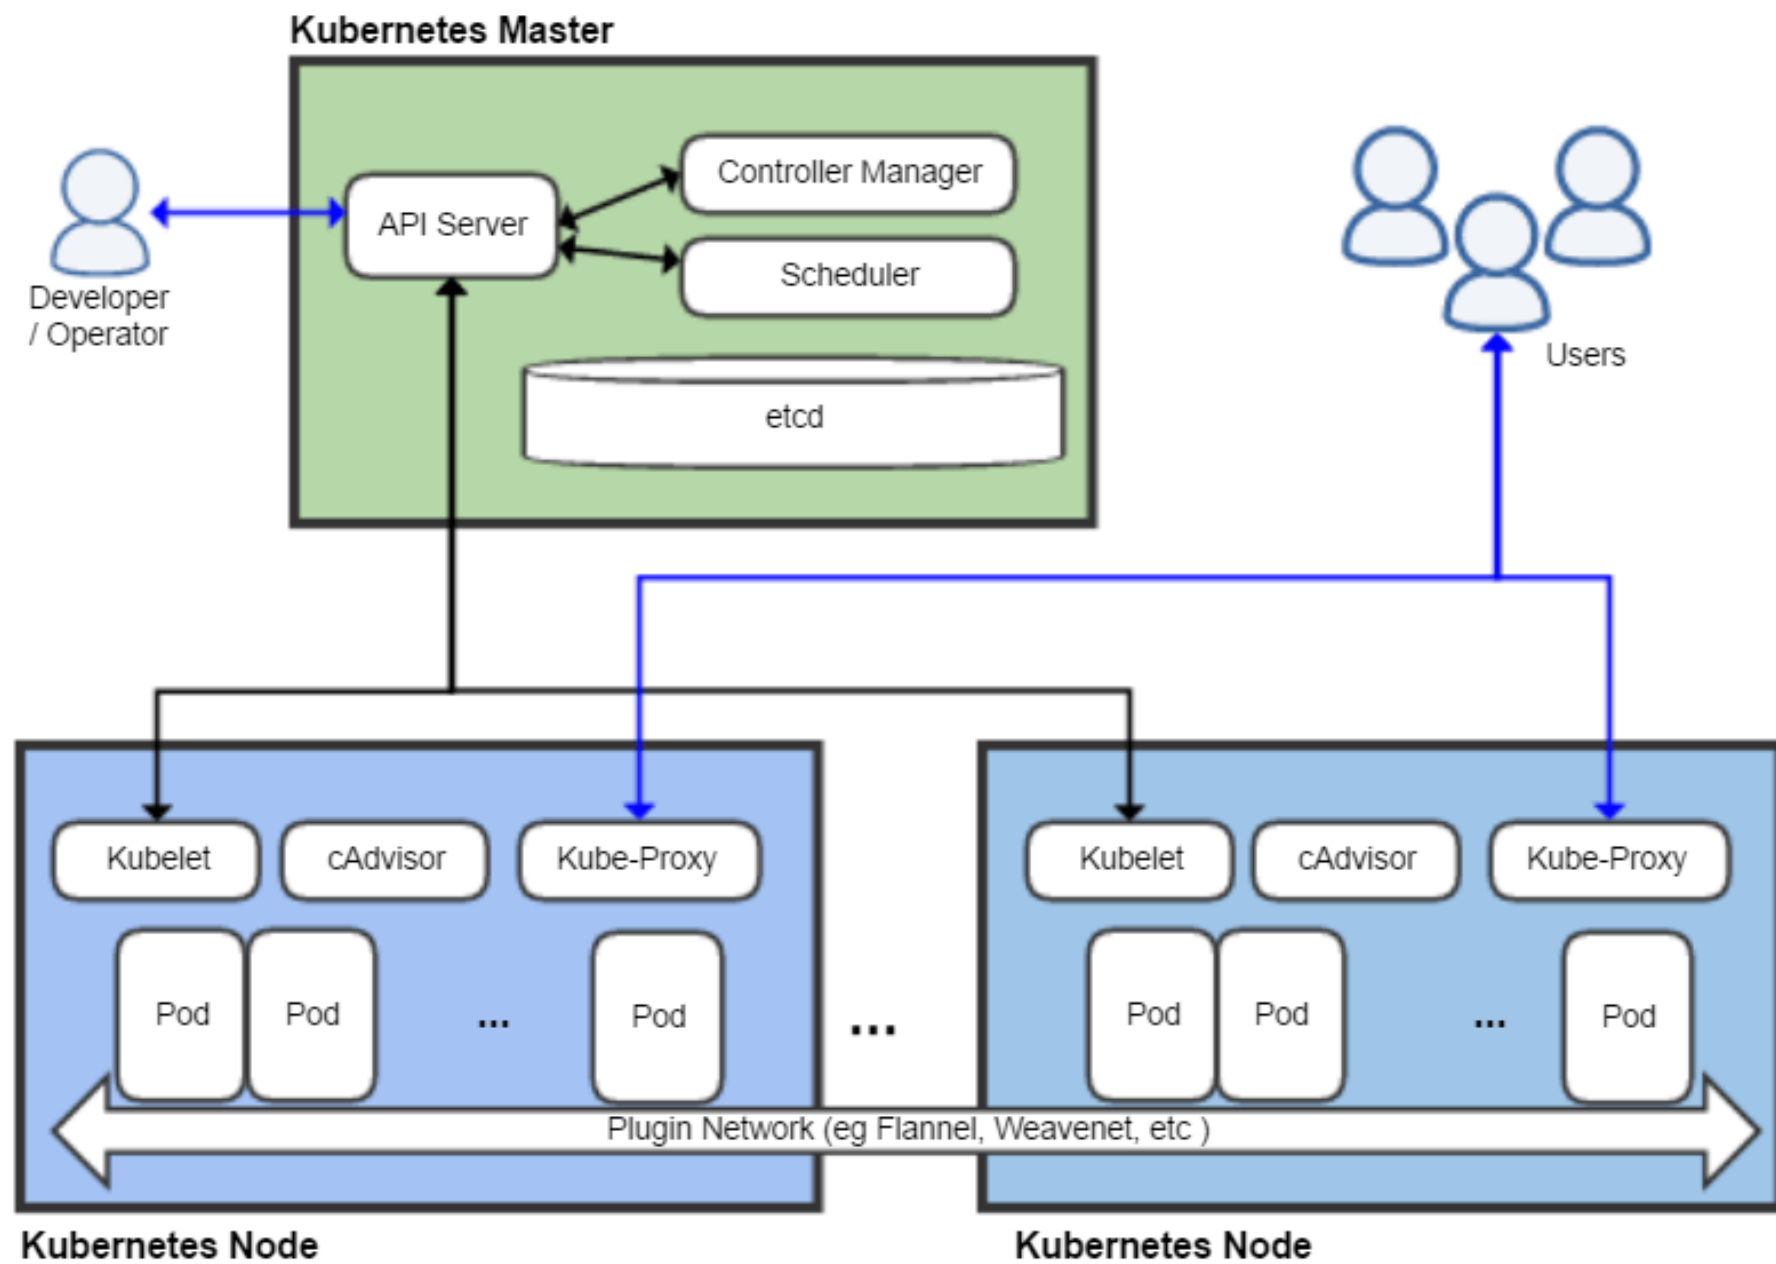
# Kubernetes

# What is Kubernetes ?

- In theory, an **orchestrator**
  - also Windows, originally, was just a **GUI** on top of DOS
- In practice, an **Operating System** for the cloud

## What is Nuvolaris?

- a Serverless **distribution** for Kubernetes
- *Linux* : **RedHat** = *Kubernetes* : **Nuvolaris**




# Kubernetes Operators

- It is a **pattern** that is becoming commonplace
  - There is *NOT* a specific API that you implement
  - You have to use the *Kubernetes API* anyway
- You define your own Resource
  - Defining new resources as **CRD** Custom Resource Definitions
  - Creating instances conforming to the CRD
    - that describes the *desired state*
  - **Writing code that brings the system to this state**

# Kubernetes

```
kubectl get nodes  
kubectl get ns  
kubectl create ns demo  
kubectl get ns
```



# Nuvolaris Architecture



# Operator Frameworks

- **Operator Framework:** ansible/helm/go
- **Kudo:** a declarative, yaml based framework
- **metacontroller:** generic, with hooks in any languages
- **shell-operator:** write operators in bash
- **kubebuilder:** Go based operator
- **kopf:** Python based Operator

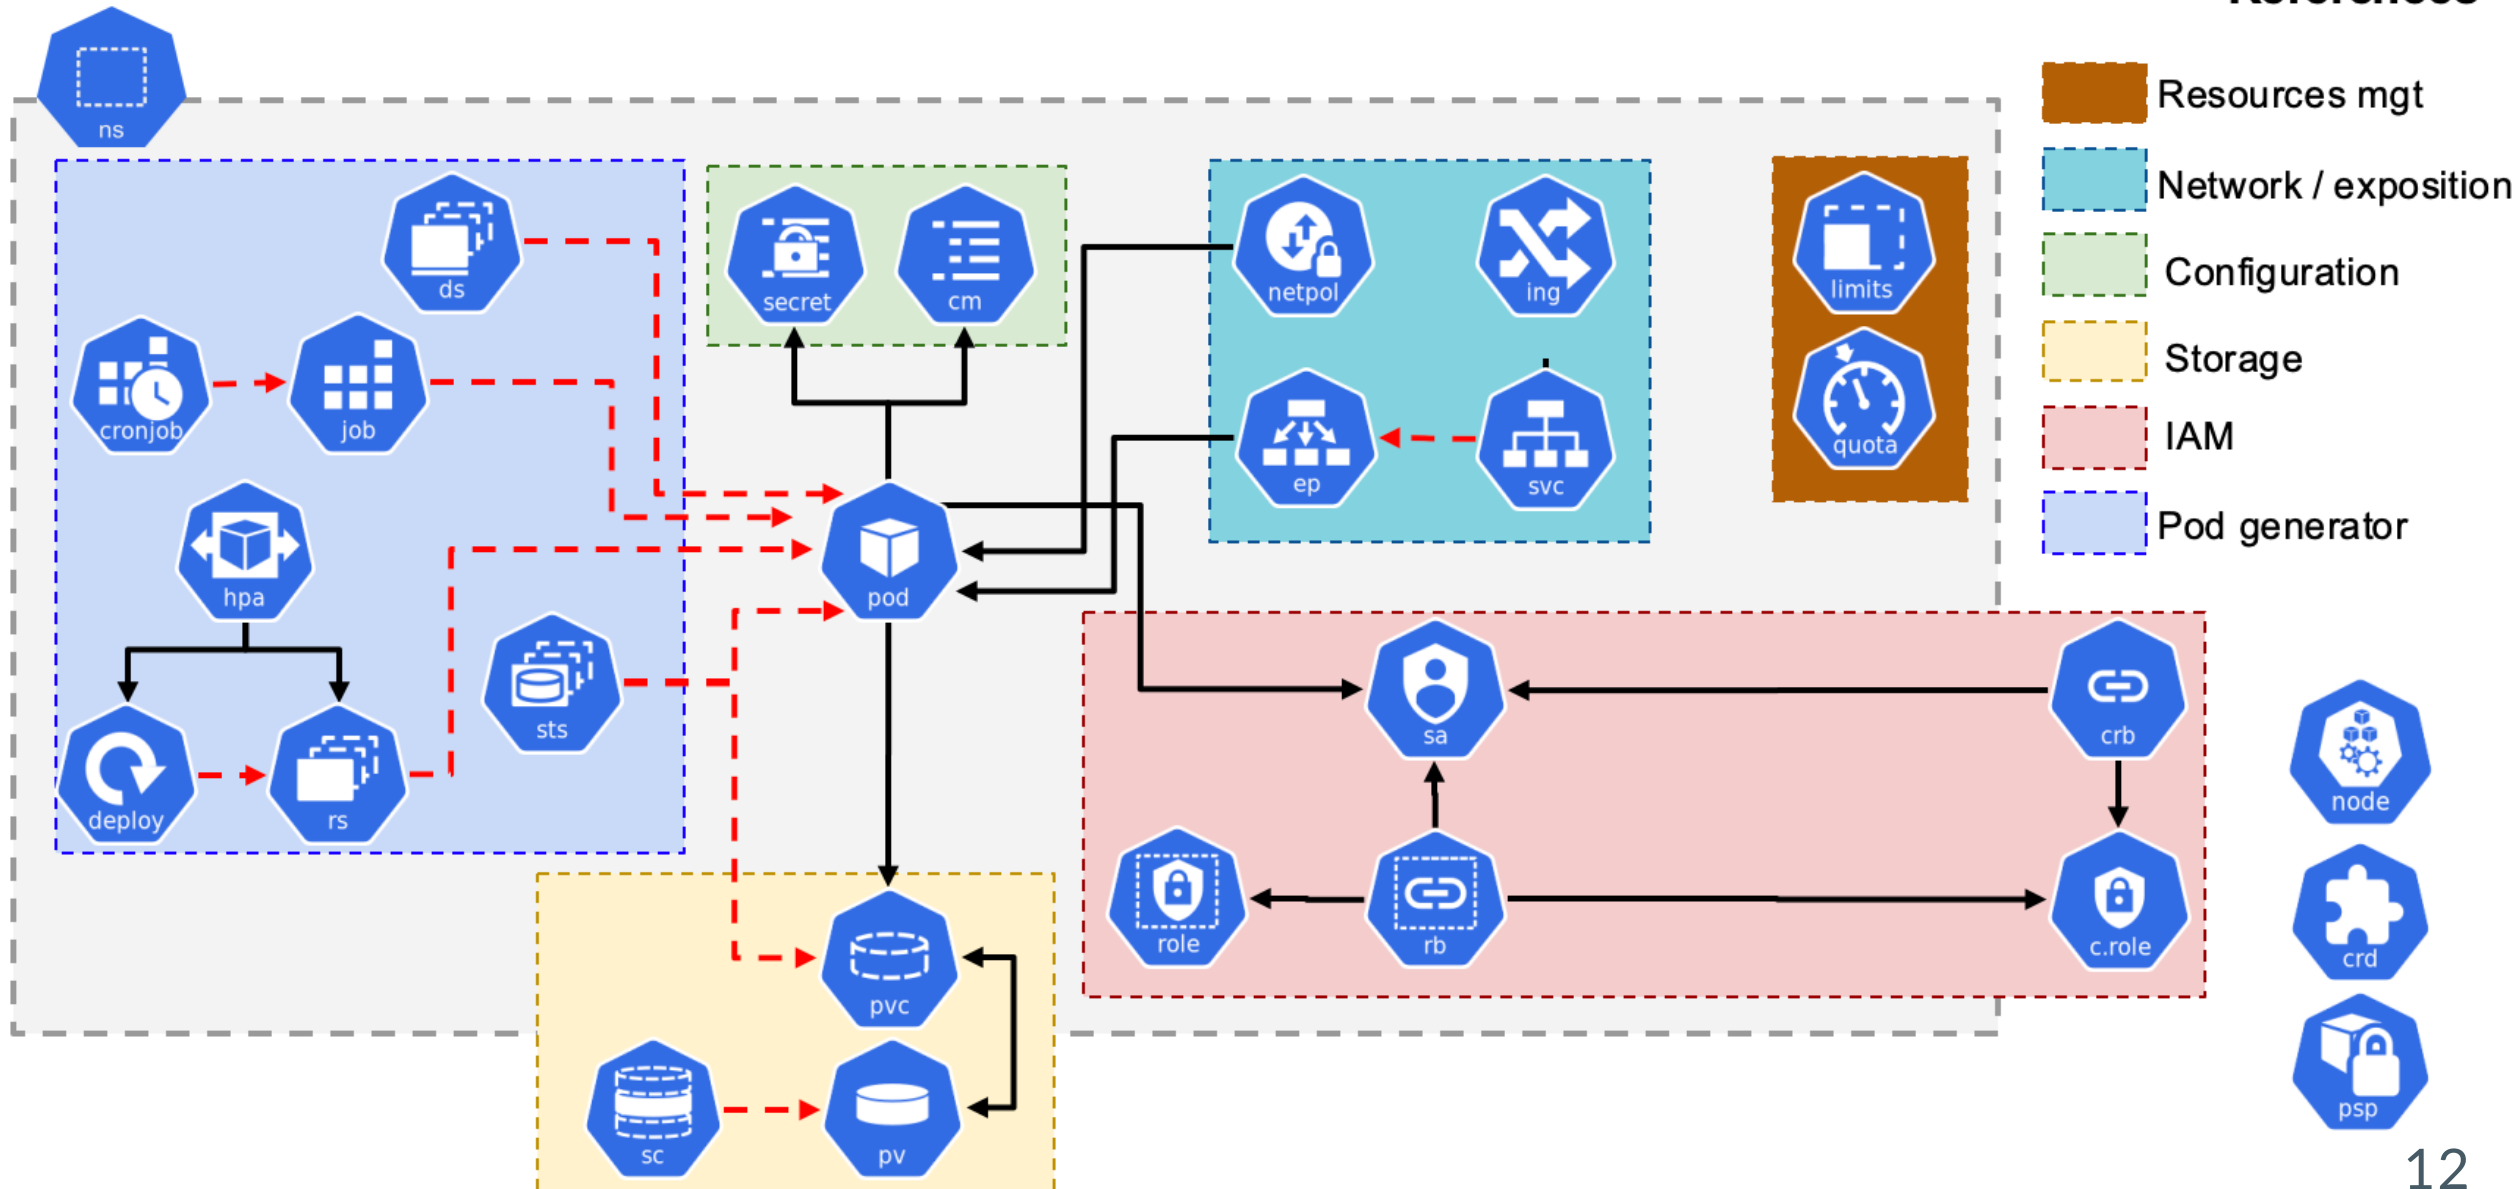
also exists Java, Rust, Elixir, Javascript based operator frameworks

# Kubernetes Descriptors

# Kubernetes Descriptors Concepts

- It is declarative:
  - You **describe** what you want to get by the system
  - Kubernetes will bring the system to the desired state
- You declare what you want with *descriptors* in YAML
  - those descriptors are in YAML format
    - actually, they are internally JSON files
    - YAML is really syntax sugar for JSON
- Kubernetes brings the system to what you asked
  - ... **if it is possible** ...

# Kubernetes Ressources Map



# Structure of a descriptor

- Common: Header and Metadata

```
apiVersion: v1
kind: Namespace
metadata:
  name: nuvolaris
```

- **spec**: changes according to the kind
- **status**: maintained by the system

# Simple Descriptor: a Pod

```
apiVersion: v1
kind: Pod
metadata:
  name: demo-pod
  namespace: demo
spec:
  containers:
    - name: nginx
      image: nginx
      ports:
        - containerPort: 80
```

# Nested Descriptor: a Deployment

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-deploy
```

Templatized, repeat the template using labels

```
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
```

# Deployment template

It creates `replica` times the pods specified in the template

```
template:
  metadata:
    labels:
      app: nginx
  spec:
    containers:
      - name: nginx
        image: nginx
        ports:
          - containerPort: 80
```



# Dev Environment

# VSCode-based Development Environment


- Clone the repositories (multiple and linked)

```
git clone https://github.com/nuvolaris/nuvolaris  
--recurse-submodules
```

- do not forget `--recurse-submodules`
- Open the folder `nuvolaris` with VSCode:
  - Command Line: `code nuvolaris`
- Open the workspaces in subfolders: `workspace.code-workspace`

1

```
$ git clone https://github.com/nuvolaris/nuvolaris --recurse-submodules  
$ code nuvolaris
```



3

workspace.code-workspace

2

Reopen in Container

# Kubernetes `kubectl`

- `$ kubectl get nodes`

NAME	STATUS	ROLES	AGE	VERSION
nuvolaris-control-plane	Ready	control-plane,master	41m	v1.21.1
nuvolaris-worker	Ready	<none>	41m	v1.21.1

**Demo:** `nuvolaris-controller/training/transcript1.txt`

# Kubernetes CRD

# Kubernetes Controllers

- Deployment, DaemonSet, StatefulSet

```
[~]$ kubectl get deploy
NAME                READY    UP-TO-DATE    AVAILABLE    AGE
nginx-deployment    3/3      3              3             21m
[~]$ kubectl get po
NAME                                READY    STATUS    RESTARTS    AGE
nginx-deployment-66b6c48dd5-4dpl2    1/1      Running    1            21m
nginx-deployment-66b6c48dd5-5c4q6    1/1      Running    1            21m
nginx-deployment-66b6c48dd5-xs8nd    1/1      Running    1            21m
```

## What they do?

- create a set of resources
- control them as an unit

# Custom Resources Definitions

- Define your own Kubernetes Resources
  - create new Kinds of resources
  - Handled as other resources

## Resource Handlers

- You need to write your own resource handler!
  - It responds to Kubernetes events
  - It interacts with Kubernetes APIs to perform operations

# Components of a CRD

- Group, Kind and short names:
  - Example: `nuvolaris.org`, `Sample`, `sam`
- Spec and Status
  - Versioned
  - defined as an OpenApi Schema:

```
type: object
properties:
  spec:
    type: object
```



# Defining a CRD (1/2)

```
apiVersion: apiextensions.k8s.io/v1
kind: CustomResourceDefinition
metadata:
  name: samples.nuvolaris.org
spec:
  scope: Namespaced
  group: nuvolaris.org
  names:
    kind: Sample
    plural: samples
    singular: sample
    shortNames:
      - sam
```

## Defining a CRD (2/2)

```
versions:
  - name: v1
    served: true
    storage: true
    subresources: { status: { } }
    schema:
      openAPIV3Schema:
        type: object
        properties:
          spec:
            type: object
            x-kubernetes-preserve-unknown-fields: true
          status:
            type: object
            x-kubernetes-preserve-unknown-fields: true
```

# Instance

```
apiVersion: nuvolaris.org/v1
kind: Sample
metadata:
  name: obj
spec:
  count: 2
```

- Demo: `nuvolaris/nuvolaris-controller/training/transcript2.txt`

# Kubernetes Operator

# About **kopf**

- See [kopf.readthedocs.io](https://kopf.readthedocs.io)
- Python based
  - provided an handy **kopf** cli runner
- Handlers for the various Kubernetes events:
  - **@kopf.on.login**
  - **@kopf.on.create**
  - **@kopf.on.delete**
- It does not manage Kubernetes API

# Login

- Kopf supports various authentication
  - Code to support either your `~/.kube/config` or the service token

```
@kopf.on.login()
def sample_login(**kwargs):
    token = '/var/run/secrets/kubernetes.io/serviceaccount/token'
    if os.path.isfile(token):
        logging.debug("found serviceaccount token: login via pykube in kubernetes")
        return kopf.login_via_pykube(**kwargs)
    logging.debug("login via client")
    return kopf.login_via_client(**kwargs)
```

# Handling object creation and deletion

```
@kopf.on.create('nuvolaris.org', 'v1', 'samples')
def sample_create(spec, **kwargs):
    print(spec)
    return { "message": "created" }
```

```
@kopf.on.delete('nuvolaris.org', 'v1', 'samples')
def sample_delete(spec, **kwargs):
    print(spec)
    return { "message": "delete" }
```

# Kustomize



# Interacting with Kubernetes

- `kopf` does *not* provide how to interact with Kubernetes
  - You can use any other api like `pykube` or others
- We use... `kubectl` and `kustomize`
  - It may look "odd" to use an external command line tool
  - However, this allows compatibility with command line tools
    - avoiding "strange" templating
    - easier development and debug

# About `kustomize`

- Originally a separate tool, now part of `kubectl`
  - It works "customizing" sets of descriptors with rules
  - support many ways of *patching* the JSON/YAML
  - **NO TEMPLATING** (huge win over `helm` !)
- You simply do `kubectl apply -k <folder>`
  - It will search for `kustomization.yaml`
  - It will produce the output sent to Kubernetes
- Debug the output without applying with:  
`kubectl kustomize <folder>`

# Simple `kustomization1.yaml` with patch

```
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
resources:
- demo-deployment.yaml
patches:
- path: patch.yaml
```

- put it in a folder `deploy` and `apply -k deploy`

# Sample patch of a Deployment

- We want to change the replica count

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-deploy
spec:
  replicas: 2
```

- Intuitively, provide enough context to locate the descriptor
- Provide the replaced fields

# Implementing Operator

# Using **kubect1** from the operator

```
# generate patch
def patch(n):
    return f"""apiVersion: apps/v1
kind: Deployment
metadata:
  name: demo-deploy
spec:
  replicas: {n}
"""

# run kubect1
def kubect1(cmd, patch):
    with open(f"deploy/patch.yaml", "w") as f:
        f.write(patch)
    res = subprocess.run(["kubect1", cmd, "-k", "deploy"], capture_output=True)
    return res.stdout.decode()
```

# Implementing the operator

```
@kopf.on.create('nuvolaris.org', 'v1', 'samples')
def sample_create(spec, **kwargs):
    count = spec["count"]
    message = kubectl("apply", patch(count))
    return { "message": message }

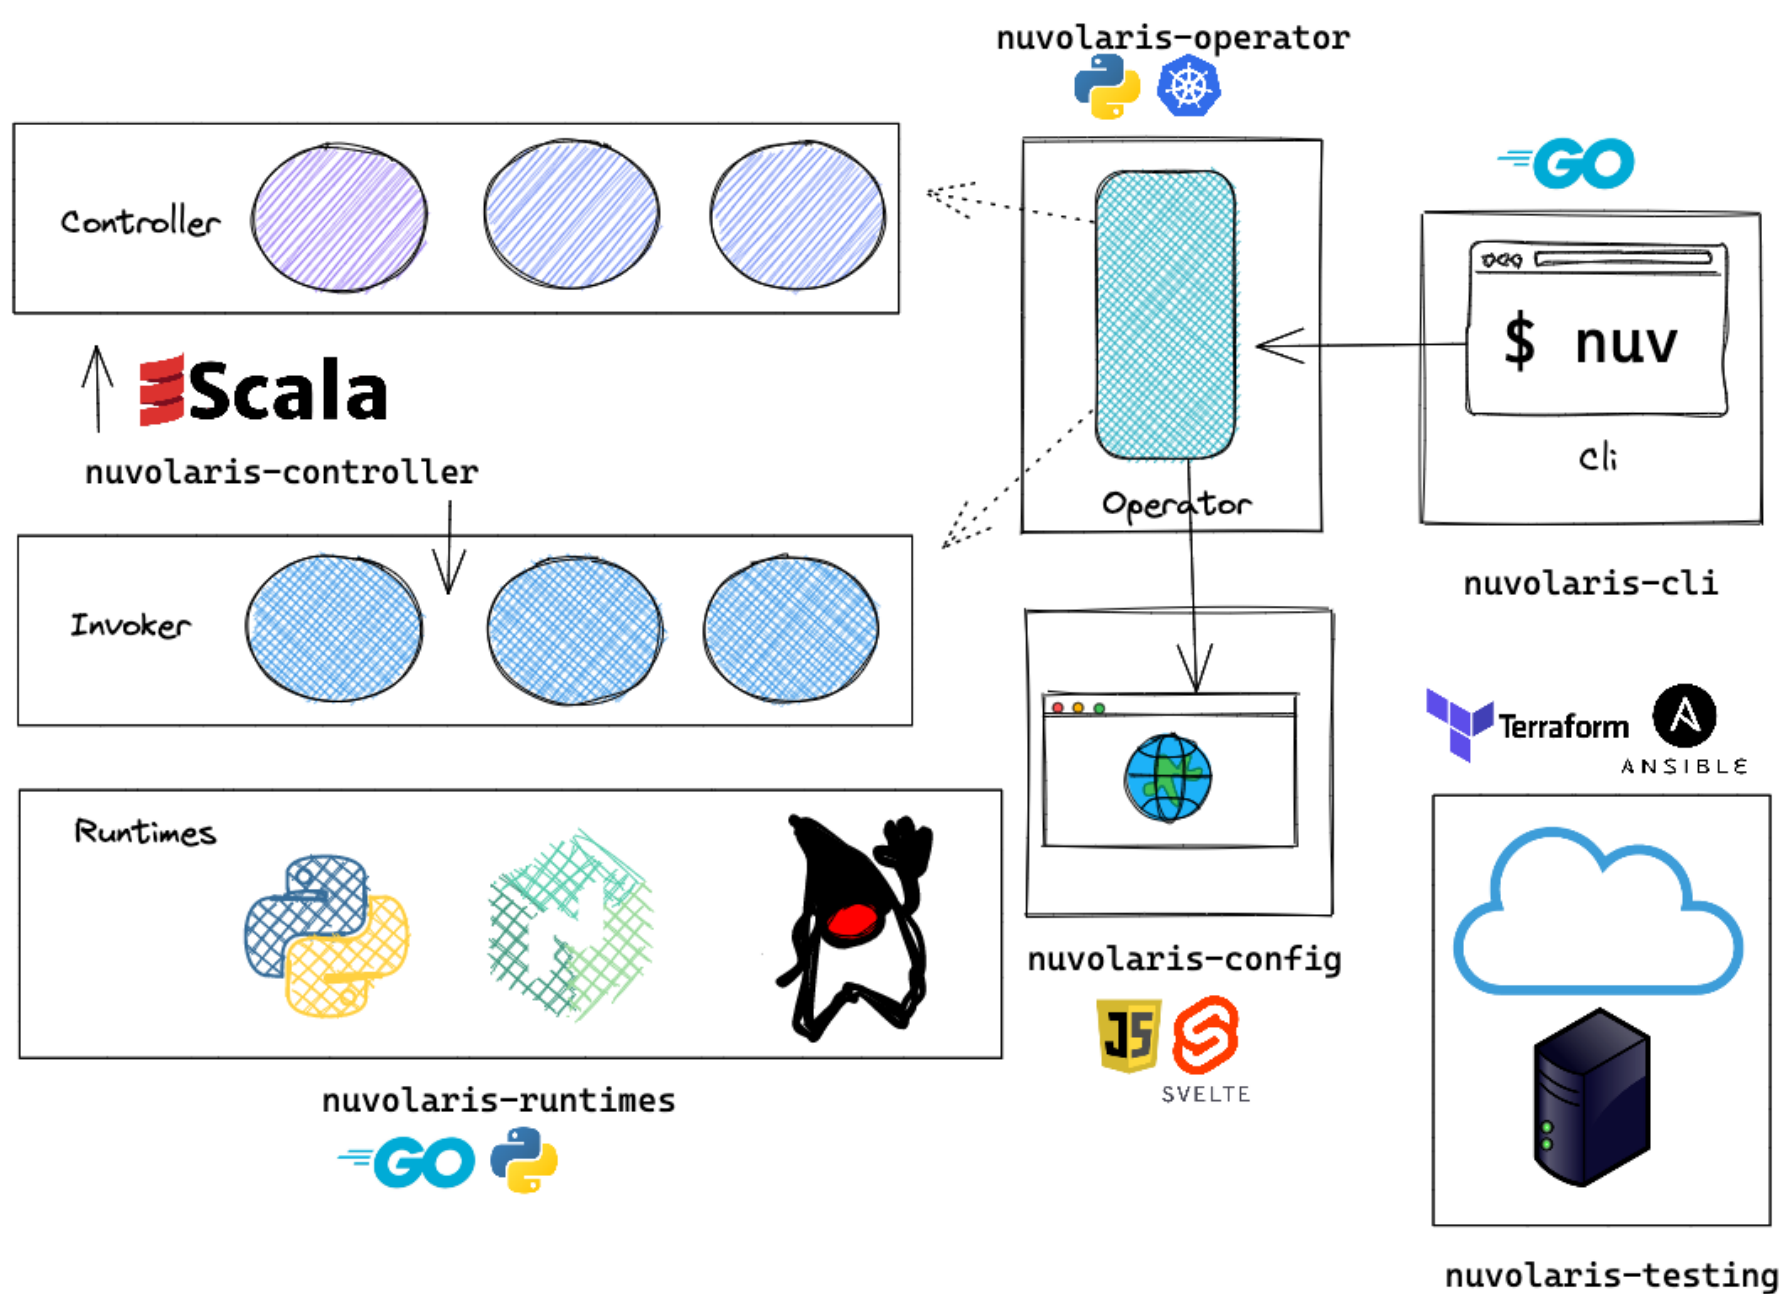
@kopf.on.delete('nuvolaris.org', 'v1', 'samples')
def sample_delete(spec, **kwargs):
    count = spec["count"]
    message = kubectl("delete", patch(count))
    return { "message": "delete" }
```

# Packaging

- Create a Dockerfile embedding the operator
  - You need `poetry`, `kopf` and `kubectl` in the image
- Deploy the POD with the right permissions
  - You need to setup *Kubernetes RBAC*
  - `ServiceAccount` and `ClusterRoleBinding`
- See `nuvolaris/nuvolaris-operator` for an example



# Contributing



## Nuvolaris Components and Technologies

# Required OpenSource Paperwork

## Before sending a PR

- Add **Apache License** headers to each file:

quick way: `license-eye header fix`

- Download the ICLA and sign it:

`bit.ly/apache-icla`

- Send to

`To: secretary@apache.org`

`Cc: secretary@nuvolaris.io`